



Data-Intensive Distributed Computing

CS 451/651 (Fall 2018)

The Final Part
November 29, 2018

Jimmy Lin
David R. Cheriton School of Computer Science
University of Waterloo

These slides are available at <http://lintool.github.io/bigdata-2018f/>



This work is licensed under a Creative Commons Attribution-Noncommercial-Share Alike 3.0 United States
See <http://creativecommons.org/licenses/by-nc-sa/3.0/us/> for details

The datacenter *is* the computer!

“Big ideas”

Scale “out”, not “up”*

Limits of SMP and large shared-memory machines

Assume that components will break*

Engineer software around hardware failures

Move processing to the data*

Clusters have limited bandwidth, code is a lot smaller

Process data sequentially, avoid random access


Seeks are expensive, disk throughput is good



Source: NASA/JPL

Humans will colonize Mars

Sooner than you think



Elon Musk Charts Path to Colonizing Mars Within a Decade

By Robin Seemangal • 06/06/16 9:10am



Planning to send Dragon to Mars by 2018. Red Dragons will inform architecture, details to come



RETWEETS
15,988 LIKES
26,640



11:50 AM - 27 Apr 2016



16K

27K



Buzz Aldrin developing a 'master plan' to colonize Mars within 25 years

Aldrin and the Florida Institute of Technology are pushing for a Mars settlement by 2039, the 70th anniversary of his own Apollo 11 moon landing



Florida Tech's president, Anthony J Catanese, left, talks with Apollo 11 astronaut Buzz Aldrin as he shows him the campus on Thursday in Melbourne, Florida. Photograph: Craig Rubadoux/AP

Source: <https://www.theguardian.com/science/2015/aug/27/buzz-aldrin-colonize-mars-within-25-years>

Source: <https://twitter.com/SpaceX/status/725351354537906176>

Source: <http://observer.com/2016/06/elon-musk-charts-path-to-colonizing-mars-within-a-decade/>

“Mars can’t just be a one-shot mission” – Buzz Aldrin



“The Pilgrims on the Mayflower came here to live and stay. They didn’t wait around Plymouth Rock for the return trip, and neither will people building up a population and a settlement [on Mars].”

Needs

Produce breathable air

Grow food

Build shelter

Mine fuel and materials

“Staying alive”

Conduct science


Connect with family and friends

Engage in leisure activities

Search the web

“Staying sane”

Maslow's hierarchy of needs



Searching the web should be as easy
from Mars as it is from Marseille!

The fundamental problem: Latency

speed of light: 2-24 minutes
rockets: 5-10 months

Bandwidth is “reasonable”

Lunar Laser Communications Demonstration:
622-Mbps downlink, 20-Mbps uplink
SneakerNet on rockets: Easily PBs

What's doable, what's not?



Example: How do I grow potatoes
in recycled organic waste?

Search from Mars: Implementation

Step 1. Rocket SneakerNet

Step 2. Beam the diffs

We know exactly how to do this!

Step 3. User model activate!

We have a good idea how to do this!

It's a caching problem!

We've worked out some simulations already...

For the truly skeptical...



Search from Mars ~
Search from regions on
Earth with poor connectivity

Easter Island
Canadian Arctic
Villages in rural India

More “down to Earth”
applications!

Big Data



What's growing faster?

Big Data

[What do I mean here?](#)

Moore's Law

[What do I mean here?](#)

Big Data > Moore's Law

Big Data < Moore's Law

Big Data ~ Moore's Law

First, a story...

What's growing faster?

Big Data

Let's restrict to
Human-generated data

Moore's Law

Bounds?

Human population
Data generation per unit time

Big Data > Moore's Law

Big Data < Moore's Law

Big Data ~ Moore's Law

Implications?

Back to my story...

X1 Instances for EC2 – Ready for Your Memory-Intensive Workloads

by Jeff Barr | on 18 MAY 2016 | in Amazon EC2* | [Permalink](#) | Share

Many AWS customers are running memory-intensive big data, caching, and analytics workloads and have been asking us for EC2 instances with ever-increasing amounts of memory.

Last fall, I first told you about our plans for the new X1 instance type. Today, we are announcing availability of this instance type with the launch of the **x1.32xlarge** instance size. This instance has the following specifications:

- Processor: 4 x Intel™ **Xeon E7 8880 v3** (Haswell) running at 2.3 GHz – 64 cores / 128 vCPUs.
- Memory: 1,952 GiB with Single Device Data Correction (SDDC+1).
- Instance Storage: 2 x 1,920 GB SSD.
- Network Bandwidth: 10 Gbps.
- Dedicated EBS Bandwidth: 10 Gbps (EBS Optimized by default at no additional cost).

The Xeon E7 processor supports **Turbo Boost 2.0** (up to 3.1 GHz), **AVX 2.0**, **AES-NI**, and the very interesting (to me, anyway) **TSX-NI** instructions. AVX 2.0 (Advanced Vector Extensions) can improve performance on HPC, database, and video processing workloads; AES-NI improves the speed of applications that make use of **AES** encryption. The new TSX-NI instructions support something cool called **transactional memory**. The instructions allow highly concurrent, multithreaded applications to make very efficient use of shared memory by reducing the amount of low-level locking and unlocking that would otherwise be needed around each memory access.

If you are ready to start using the X1 instances in the US East (Amazon Pacific (Singapore), or Asia Pacific (Sydney) Regions, please [request](#) instances available in other Regions and in other sizes before t

```
top - 17:13:41 up 52 min, 3 users, load average: 10.52, 4.74, 2.70
Tasks: 943 total, 2 running, 941 sleeping, 0 stopped, 0 zombie
%Cpu(s): 0.2 us, 0.6 sy, 0.0 ni, 94.4 id, 4.8 wa, 0.0 hi, 0.0 si, 0.
KiB Mem: 20184757+total, 87502960 used, 19309728+free, 19360 buffers
KiB Swap: 0 total, 0 used, 0 free. 40866816 cached Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+
35957	hdbadm	20	0	86.849g	0.065t	0.036t	S	118.812	3.458	8:38.77
35959	hdbadm	20	0	5643852	1.491g	287768	S	1.650	0.077	0:32.41
1120	root	20	0	12032	3792	664	S	0.990	0.000	0:02.69
37350	root	20	0	14472	2228	1008	R	0.660	0.000	0:00.20
8	root	rt	0	0	0	0	S	0.330	0.000	0:00.53

What's growing faster?

Big Data

Let's restrict to
Human-generated data


Moore's Law


What am I forgetting?

Bounds?


Human population
Data generation per unit time

Serverless Architectures







Server




Server




Server




Server




Server



Cloud (I'm going to illustrate with AWS)




Cloud




Cloud


“State” as a service (S3, RDS, SQS, ...)




Cloud



Cloud



Cloud



Serverless Architectures


Doesn't mean you don't have servers

Just that managing them is the cloud provider's problem


Write functions with well-defined entry and exit points

Cloud provider handles all other aspect of execution

Example: Image Thumbnail Creation



Example: Analysis of Streaming Social Media Data



(Current) Serverless Architectures


Asynchronous, loosely-coupled, event-driven
Functions touch relatively little data

What about serverless data analytics?

Design goal: pure pay-as-you-go, zero costs for idle capacity
Compared to current options?

Flint

PySpark execution backend



The datacenter *is* the computer!

“Big ideas”

Scale “out”, not “up”*

Limits of SMP and large shared-memory machines

Assume that components will break*

Engineer software around hardware failures

Move processing to the data*

Clusters have limited bandwidth, code is a lot smaller

Process data sequentially, avoid random access

Seeks are expensive, disk throughput is good

